

SQUIC

SOFTWARE FOR SPARSE PRECISION MATRIX ESTIMATION

User Manual

Version 1.0

November 20, 2023

Table of Contents

1	Introduction	2
2	Compiling and Installation	3
3	SQUIC for R	4
3.1	List of variables	4
3.2	Usage	5
4	SQUIC for Python	7
4.1	List of variables	7
4.2	Usage	8
5	Example Case Study	10

About this document

This file is intended as a user manual for the software package SQUIC. The package and its source code are available at www.gitlab.ci.inf.usi.ch/SQUIC.

1 Introduction

This software is an implementation of the **S**parse **Q**Uadratic approximation of **I**nverse **C**ovariance matrix (SQUIC) algorithm [3], a second-order method designed for the solving ℓ_1 -regularized Gaussian maximum likelihood problem. Based on the foundational work of [5], this algorithm is developed for large-scale, and performance-oriented sparse inverse covariance or precision matrix estimation. This user manual serves as a resource for the SQUIC package and its interfaces.

For the dataset $\mathbf{Y} \in \mathbb{R}^{p \times n}$, consisting of p random variables and n samples, SQUIC computes a sparse estimate of the precision matrix and its sparse approximate inverse, denoted as $\hat{\Theta}$ and $\hat{\Theta}^{\text{inv}}$, respectively. The estimation process employs a *matrix tuning parameter* $\Lambda \in \mathbb{R}^{p \times p}$, which is indirectly defined using $\mathbf{M} \in \mathbb{R}^{p \times p}$ and a *scalar tuning parameter* $\lambda > 0$. This relationship is summarized as follows:

$$\Lambda_{ij} := \begin{cases} \mathbf{M}_{ij}, & \text{if } \mathbf{M}_{ij} \neq 0, \\ \lambda & \text{otherwise.} \end{cases} \quad (1)$$

In many cases, it is appropriate to define the nonzero entries of $\mathbf{M}$ as constants, such that, for example, for $\eta \in (0, \lambda]$, $\mathbf{M}_{ij} = \eta$ for all its nonzero entries. The larger the value of Λ_{ij} , the more the corresponding value $\hat{\Theta}_{ij}$ is penalized for being nonzero. If $\mathbf{M}$ is not specified and only the scalar tuning parameter λ is specified, we set $\Lambda_{ij} = \lambda$ for all its entries. In Figure 1, we visualize the main inputs and outputs of SQUIC, alongside its primary computational components.

The source code for the software described in this manuscript is available under general public licenses (GPL 3) at www.gitlab.ci.inf.usi.ch/SQUIC. The software is packaged as `libSQUIC` and intended for use on Linux and Mac systems (both X86 and ARM architectures). It is written in C++, parallelized with OpenMP, and features interface packages for R and Python. This codebase incorporates the SQUIC algorithm from [1], along with parallel and blocking components in [2] and [4], respectively.




Figure 1: A visualization of the SQUIC package which includes the main inputs $\mathbf{Y}$ and Λ , as well as the outputs $\hat{\Theta}$ and $\hat{\Theta}^{\text{inv}}$, corresponding to the data matrix, matrix tuning parameter, estimated precision matrix and its sparse approximate inverse, respectively. Note that we can set $\Lambda_{ij} = \lambda$ for all its entries, where λ is the scalar tuning parameter; see (1) for details. Also shown are the main computational components, which include the upfront sparse sample covariance matrix computation (blue), alongside the Newton iteration (red), which includes block coordinating descent, supernodal Cholesky factorization, and block approximate matrix inversion. Shared-memory parallel components are marked with an “*”.

2 Compiling and Installation

The SQUIC library, called `libSQUIC` is hosted at www.gitlab.ci.inf.usi.ch/SQUIC/lib. In this section, we outline two methods for installing the library: using available precompiled versions or compiling the library from source. For both the R and Python interfaces `libSQUIC` is required. The default and recommended location of the library is `HOME`.

Precompiled: This can be downloaded directly from the repository in `/precompiled` folder, where different precompiled versions of `libSQUIC` can be found. Depending on the version, there may be a dependency on OpenMP. For Mac, OpenMP can be installed with `$ brew install libomp`.

From Source: Compiling `libSQUIC` from source requires the following packages:

- OpenMP
- Compilers C/C++/Fortran

For Mac users, using the Homebrew package manager is recommended. You can install the GNU Fortran compiler, and OpenMP with Homebrew by executing `$ brew install gcc libomp` (the default Mac Clang compiler suffices for C/C++). With the required packages installed, the following steps will compile the source code:

- ❶ Clone the repository.

```
$ git clone https://www.gitlab.ci.inf.usi.ch/SQUIC/lib SQUIC/lib
```

- ❷ Compile the dependencies (refer to supporting libraries for CHOLMOD in `extern/src`) using `$ make dependencies`. After compiling the dependencies, compile the SQUIC library with `$ make squic`. Both steps can be executed by just running `$ make`. Finally, install the shared library in the default `HOME` location using `$ make install`. The entire compilation and installation process can be performed with the following commands:

```
$ make
$ make install
```

- ❸ To run a basic demonstration of the Python interface for the `libSQUIC` library and its Python interface package, use this command:

```
$ make demo
```

To compile `libSQUIC` with Intel Math Kernel Library (MKL), set the `MKLROOT` environment variable by executing `$ source /opt/intel/oneapi/mkl/latest/env/vars.sh intel64` (note: this command may vary based on your MKL version). Then, adjust your compiler and linker flags according to your system. For specific flag details, refer to the [Intel MKL Link Line Advisor](#).

3 SQUIC for R

The R interface for SQUIC is available at www.gitlab.ci.inf.usi.ch/SQUIC/R. You can install the package by following these steps:

- ❶ Install `libSQUIC` in the recommended home directory, i.e., `HOME`. See Section 2 for details on precompiled versions, and compiling from source.
- ❷ Install the library with the following R commands.

```
library(devtools)
install_git("https://www.gitlab.ci.inf.usi.ch/SQUIC/SQUIC_R.git")
```

The environment variable `SQUIC_LIB_PATH` defines the location of `libSQUIC`. If this variable is not set, the default location of `libSQUIC` is `HOME`. To change the default location, in R use the command `Sys.setenv(SQUIC_LIB_PATH="/path/to/squic/")`, replacing `/some/to/squic` with the path of where `libSQUIC` can be found.

3.1 List of variables

The `SQUIC()` function call serves as the interface function for `libSQUIC`, invoking the SQUIC algorithm for sparse precision matrix estimation. The list of input and output variables is shown below and can also be viewed via the help command `help(SQUIC)`.

```
out1 <- SQUIC(Y,lambda,max_iter,tol,tol_inv,verbose,M,X0,W0);
out2 <- SQUIC(Y,lambda,max_iter = 0,tol,verbose,M)
```

All sparse matrix formats utilize the standard `Matrix` package.

Inputs

- `Y`: Input data matrix ($p \times n$ dense matrix).
- `lambda`: Scalar tuning parameter a nonzero positive number.
- `max_iter` [default 100]: Maximum number of Newton iterations¹
- `tol` [default 10^{-3}]: Tolerance for convergence.
- `tol_inv` [default 10^{-4}]: Tolerance for approximate matrix inversion.
- `verbose` [default 1]: Verbosity level as 0 or 1.
- `M` [default `NULL`]: Symmetric sparse matrix ($p \times p$ sparse matrix).²
- `X0` [default `NULL`]: Initial value of precision matrix ($p \times p$ sparse matrix).³
- `W0` [default `NULL`]: Initial value of inverse precision matrix ($p \times p$ sparse matrix).³

¹Setting `max_iter=0` will return the *sparse (thresholded) sample covariance matrix*.

² We define $\Lambda_{ij} = M_{ij}$ if $M_{ij} \neq 0$ and $\Lambda_{ij} = \lambda$, if otherwise. When `M=NULL` (or `None` in Python), we default to the, so called, “scalar-tuning parameter” for which $\Lambda_{ij} = \lambda$.

³If either `X0` or `W0` is set to `NULL` (or `None` in Python), both `X0` and `W0` are set to identity.

Outputs

- `out1$X` : Estimated precision matrix ($p \times p$ sparse matrix).
- `out1$W` : Sparse approximate inverse of the precision matrix ($p \times p$ sparse matrix).
- `out1$info_time_total` : Total runtime.
- `out1$info_time_sample_cov` : Runtime of the sample covariance matrix.
- `out1$info_time_optimize` : Runtime of the Newton steps.
- `out1$info_time_factor` : Runtime of matrix factorization.
- `out1$info_time_approximate_inv` : Runtime of approximate matrix inversion.
- `out1$info_time_coordinate_upd` : Runtime of coordinate descent update.
- `out1$info_objective` : Objective value at each Newton iteration.
- `out1$info_logdetX` : Log determinant of the estimated precision matrix.
- `out1$info_trSX` : Trace of sample cov. times the estimated precision matrix.

- `out2$S` : Sparse sample covariance matrix ($p \times p$ sparse matrix).
- `out2$info_time_total` : Total runtime.
- `out2$info_time_sample_cov` : Runtime of the sample covariance matrix.

3.2 Usage

In the basic example outlined below, we will use SQUIC to estimate the precision matrix of a synthetically generated dataset with correlated random variables, where the true precision matrix is tridiagonal. For an interactive session of the R SQUIC interface see [Google Colabs](#).⁴

```
library(SQUIC)

# basic test parameters
p      <- 1024
n      <- 100
lambda <- .4

# generate a tridiagonal matrix & dataset
set.seed(1)
iC_star <- Matrix::bandSparse(p, p, (-1):1, list(rep(-.5, p-1), rep(1.25, p),
  ~ rep(-.5, p-1)));
z      <- replicate(n, rnorm(p))
iC_L   <- chol(iC_star)
Y      <- matrix(solve(iC_L, z), p, n)

# run SQUIC
out <- SQUIC(Y, lambda)
```

⁴ There may be slight variations in the results across different machines due to *round-off errors*. These differences are visible in the sparse approximation of the estimated inverse of the precision matrix `W` and are associated with the specific permutation used by the underlying matrix factorization routine, which are not related to the random seed for data generation.

With the default verbosity level the output is as shown below.⁵

```
-----
                        SQUIC Version 1.0
-----

Input Matrices
nnz(X0)/p:  1.000000e+00
nnz(W0)/p:  1.000000e+00
nnz(M)/p:   ignored
Y:          1024 x 100

Runtime Configs
Sample Cov.:  Deterministic
CordDec Vers: Collective coordinate descent update
Inversion:    Approx. block Neumann series
Fact. Routine: CHOLMOD

Parameters
verbose:      1
lambda:       4.000000e-01
max_iter:     100
term_tol:     1.000000e-03
inv_tol:      1.000000e-04
threads:      12

#SQUIC Started
* sample covariance matrix S: time=2.36e-03 nnz(S)/p=6.91e+00
* iter=1 time=4.89e-03 obj=1.60e+03 |delta(obj)|/obj=1.56e-01 nnz(X,L,W)/p=
  [6.91e+00 7.54e+01 1.69e+00] lns_iter=3
* iter=2 time=7.23e-03 obj=1.56e+03 |delta(obj)|/obj=3.13e-02 nnz(X,L,W)/p=
  [6.91e+00 7.54e+01 5.75e+01] lns_iter=2
* iter=3 time=3.19e-02 obj=1.55e+03 |delta(obj)|/obj=6.62e-03 nnz(X,L,W)/p=
  [6.91e+00 7.54e+01 5.78e+01] lns_iter=1
* iter=4 time=3.06e-02 obj=1.54e+03 |delta(obj)|/obj=2.71e-04 nnz(X,L,W)/p=
  [6.91e+00 7.54e+01 9.01e+01] lns_iter=1
#SQUIC Finished: time=8.60e-02 nnz(X,W)/p=[6.91e+00 9.01e+01]
```

Input details are provided under “Input Matrices”, where X_0 and W_0 denote the initial values of the precision matrix and its inverse (defaulting to identity matrices). The absence of M is noted, and the size of the input data Y is specified as $p = 1024$, $n = 100$. In the “Runtime Configs” section, key computational operations are summarized. The “Parameters” section lists parameters used in the overall optimization procedure, where the maximum number of threads available on the system is automatically utilized. Here `iter`, `nnz`, and `lns_iter` represent the Newton iteration index, number of nonzeros, and the number of line-search iterations, respectively. For instance, `nnz(X,L,W)/p` indicates the number of nonzeros per row for the precision matrix, its Cholesky factor, and the sparse approximate inverse, respectively.

⁵In R notebooks the verbose output may not be visible from within the notebook (see, e.g., [here](#)). This issue is not present when using the terminal or RStudio, and it does not affect the results.

4 SQUIC for Python

The Python (V3) interface for SQUIC is available at www.gitlab.ci.inf.usi.ch/SQUIC/python. You can install the package by following these steps:

- ❶ Install `libSQUIC` in the recommended and default directory `HOME`. For details on precompiled versions, and compiling from source, see Section 2.
- ❷ Install the interface package using the Python package manager.

```
$ pip3 install squic
```

Use `$ export SQUIC_LIB_PATH=/path/to/squic/`, replacing `/path/to/squic/` with the path of where `libSQUIC` can be found. If this variable is not set, the default location of `libSQUIC` is `HOME`. This should be done before calling Python.

4.1 List of variables

The `squic.run()` function call serves as the Python interface function for the `libSQUIC`, invoking the SQUIC algorithm for sparse precision matrix estimation. The list of input and output variables is shown below and can also be viewed via the help command `help(squic)`. The description of the function here is exactly the same as the R interface discussed in Section 3. The sole difference between the input variable of the R and Python interface is that `lambda` is a reserved keyword in Python, so `l` is used instead. For completeness, the input and output variable descriptions are provided below.

```
out1 = squic.run(Y,l,max_iter,tol,tol_inv,verbose,M,X0,W0);
out2 = squic.run(Y,l,max_iter = 0,tol,verbose,M);
```

All sparse matrix formats utilize the `scipy` package.

Inputs

- `Y`: Input data matrix ($p \times n$ dense matrix).
- `l`: Scalar tuning parameter a nonzero positive number.
- `max_iter` [default 100]: Maximum number of Newton iterations.¹
- `tol` [default 10^{-3}]: Tolerance for convergence.
- `tol_inv` [default 10^{-4}]: Tolerance for approximate matrix inversion.
- `verbose` [default 1]: Verbosity level as 0 or 1.
- `M` [default `None`]: Symmetric sparse matrix ($p \times p$ sparse matrix)..²
- `X0` [default `None`]: Initial value of precision matrix ($p \times p$ sparse matrix).³
- `W0` [default `None`]: Initial value of inverse precision matrix ($p \times p$ sparse matrix).³

Outputs

- `out1[0]` : Estimated precision matrix ($p \times p$ sparse matrix) .
- `out1[1]` : Sparse approximate inverse of the precision matrix ($p \times p$ sparse matrix).
- `out1[2]` : List of different component runtimes.
 - `out1[2][0]` : Total runtime.
 - `out1[2][1]` : Runtime of the sample covariance matrix.
 - `out1[2][2]` : Runtime of the Newton steps.
 - `out1[2][3]` : Runtime of matrix factorization.
 - `out1[2][4]` : Runtime of approximate matrix inversion.
 - `out1[2][5]` : Runtime of coordinate descent update.
- `out1[3]` : Objective value at each Newton iteration.
- `out1[4]` : Log determinant of the estimated precision matrix..
- `out1[5]` : Trace of sample covariance matrix times the estimated precision matrix.
- `out2[1]` : Sparse sample covariance matrix ($p \times p$ sparse matrix) .
- `out2[2]` : List of different compute times.
 - `out2[2][0]` : Total runtime.
 - `out2[2][1]` : Runtime of the sample covariance matrix.

4.2 Usage

In the basic example outlined below, we will use SQUIC to estimate the precision matrix of a synthetically generated dataset with correlated random variables, where the true precision matrix is tridiagonal. For an interactive session of the Python SQUIC interface see [Google Colabs](#).⁴

```
import squic
import numpy as np

# basic test parameters
p = 1024
n = 100
l = .4

# generate a tridiagonal matrix & dataset
np.random.seed(1)
a = -0.5 * np.ones(p-1)
b = 1.25 * np.ones(p)
iC_star = np.diag(a,-1) + np.diag(b,0) + np.diag(a,1)
L = np.linalg.cholesky(iC_star)
Y = np.linalg.solve(L.T,np.random.randn(p,n))

# run SQUIC
[X,W,info_times,info_objective,info_logdetX,info_trSX] = squic.run(Y,l)
```

Using the default verbosity level the output is as shown below. This output differs compared to the R example in Section 3, as the generated random data differs. With this said, reruns of the same script, using the same random seed, will result in the same results.

```
-----
                        SQUIC Version 1.0
-----

Input Matrices
nnz(X0)/p:  1.000000e+00
nnz(W0)/p:  1.000000e+00
nnz(M)/p:   ignored
Y:          1024 x 100

Runtime Configs
Sample Cov.:  Deterministic
CordDec Vers: Collective coordinate descent update
Inversion:    Approx. block Neumann series
Fact. Routine: CHOLMOD

Parameters
verbose:      1
lambda:       4.000000e-01
max_iter:     100
term_tol:     1.000000e-03
inv_tol:      1.000000e-04
threads:      12

#SQUIC Started
* sample covariance matrix S: time=1.91e-02 nnz(S)/p=6.49e+00
* iter=1 time=4.34e-02 obj=1.60e+03 |delta(obj)|/obj=1.60e-01 nnz(X,L,W)/p=
  [6.49e+00 6.61e+01 1.67e+00] lns_iter=3
* iter=2 time=6.63e-02 obj=1.55e+03 |delta(obj)|/obj=3.11e-02 nnz(X,L,W)/p=
  [6.49e+00 6.61e+01 5.30e+01] lns_iter=2
* iter=3 time=4.80e-02 obj=1.54e+03 |delta(obj)|/obj=6.42e-03 nnz(X,L,W)/p=
  [6.49e+00 6.61e+01 5.29e+01] lns_iter=1
* iter=4 time=4.44e-02 obj=1.54e+03 |delta(obj)|/obj=2.53e-04 nnz(X,L,W)/p=
  [6.49e+00 6.61e+01 8.24e+01] lns_iter=1
#SQUIC Finished: time=2.39e-01 nnz(X,W)/p=[6.49e+00 8.24e+01]
```

Note that the verbose output of `libSQUIC` is not visible in notebooks; but it does not affect the results.⁶ This issue does not occur when using Python in a non-notebook environment. By using the `wurlitzer` package, we can sidestep this issue with the following lines of code before the `squic.run` notebook function call.

```
!pip install wurlitzer
%load_ext wurlitzer
```

⁶This is a known issue; see, e.g., pypi.org/project/wurlitzer for details.

5 Example Case Study

In this section, we present a case study that employs SQUIC for classifying high-dimensional gene expressions through linear discriminant analysis (LDA). The complete R source code for this experiment is available at www.gitlab.ci.inf.usi.ch/SQUIC/replication. This section outlines the key computational steps of the experiment in the form of R “code snippets.” The description presented here is intended to complement the text in Section 4.2 of [3], where the results, including the classification accuracy and timings of the conducted tests are outlined and discussed in detail. The collection of datasets used for our test are available from the `datamicroarray` package.⁷ In the outlined code snippets, we use the transcription of the variable name; for example, η or $\mathbf{M}$ are denoted as `eta` and `M`, respectively.

We begin by loading the datasets and employ a stratified approach to split them into training and testing sets. Using the `caret` package,⁸ we select 70% of the samples, comprising of the feature values and their corresponding labels from each class, to form the training set. This data is referred to as `data_train` and `labels_train`. The remainder of the dataset is used for testing, for which we follow the same naming standard. To simplify the selection of the tuning parameter, matrix or its scalar form (see, e.g., Section 1 for further details), we normalized the data using the diagonal matrix `D` as shown below. We later undo this scaling once the precision matrix estimation procedure is completed.

```
n_train <- dim(data_train)[2]
a <- (apply(data_train,1,var))*(n_train-1)/n_train
D <- Diagonal(x=1/sqrt(a))
data_train_scaled <- D%*%data_train
```

In this study, we define two versions of the tests: one using the matrix tuning parameter $\mathbf{\Lambda}$, and the other using just the scalar tuning parameter λ . To define $\mathbf{\Lambda}$, we need to first specify $\mathbf{M}$, which we now describe. In the first step, we generate a complete graph with edge weights represented by the Euclidean distance. Subsequently, we compute the adjacency matrix, denote here as `G`, corresponding to the Minimal Spanning Tree (MST). The MST is a sparsely connected graph with the minimum sum of edge weights required for connectivity. The nonzero pattern of `G` corresponds to pairs of nodes with low Euclidean distances between them, indicating high similarity. By using the `emstreeR` package,⁹ `G` can be efficiently computed as follows.

```
mst <- ComputeMST(as.matrix(data_train))
G <- sparseMatrix(dims=c(p,p), i=mst$to, j=mst$from, x=rep(1, p))
G <- forceSymmetric(G,uplo="L") + Diagonal(nrow(G))
```

⁷The `datamicroarray` package offers a collection of microarray datasets to assess machine learning algorithms and is available at github.com/ramhiser/datamicroarray.

⁸The `caret` package is a collection of functions for training and plotting classification and regression models, and is available at cran.r-project.org/package=caret.

⁹The `emstreeR` package computes an Euclidean MST from data, and is available at cran.r-project.org/package=emstreeR.

The last operation enforces symmetry and adds diagonal entries to $\mathbf{G}$, as the adjacency matrix of the MST will not have diagonal entries or self-links. Note that the diagonal entries of $\mathbf{G}$ are added simply so that they are reflected in $\mathbf{M}$ —the adjacency matrix of a MST will have zero for its diagonal entries. We can now define $\mathbf{M} \leftarrow \mathbf{G} * \eta$, where for this study we fix $\eta \leftarrow 0.1$.

Using SQUIC, we now compute the required precision matrices, for which an appropriate λ must be selected. In this example, we set $\lambda = 0.8$, but this choice depends on the dataset. For further details on the selection of λ for the different datasets, please refer to [3] or the replication script. For the estimation procedure using the matrix tuning parameter, we fix $\lambda = 0.95$ across all dataset tests. The respective function calls for using just the scalar tuning parameter, and the case where the matrix tuning parameter is used, are as follows:

```
squic_out <- SQUIC(Y=data_train,lambda=0.80,M=NULL,tol=1e-3,tol_inv=1e-3)
squic_out_M <- SQUIC(Y=data_train,lambda=0.95,M=eta*G,tol=1e-3,tol_inv=1e-3)
```

After computing $\hat{\Theta}$, we *undo* the normalization operation; for example, for the scalar tuning estimate, we would rescale the matrix as follows: $\mathbf{D} \% \% \text{squic_out}\$X \% \% \mathbf{D}$.

The LDA method is subsequently applied to classify the samples in the testing dataset. For each class $k = 1, \dots, K$, where K is the number of classes in the given dataset (i.e., the labels), we compute the sample means using `data_train`. Denoting the estimated precision matrices as `theta_est`, we compute the LDA score, `lda_score`, for each class and assign a predicted class based on the one with the highest LDA score.

```
for (k in 1:K) {
  # prior probability for class k
  prior <- length(which(labels_train == k)) / n_train
  # sample mean for class k
  mu <- rowMeans(data_train[, which(labels_train == k)])
  mu_matrix <- matrix(mu, nrow = nrow(data_test), ncol = ncol(data_test), byrow =
    FALSE)
  # x - 0.5*\mu
  adjusted_data <- data_test - 0.5 * mu_matrix
  # Linear Discriminant function
  lda_score <- diag(t(adjusted_data) \% \% theta_est \% \% mu_matrix + log(prior))
  rho[, k] <- lda_score
}
predicted <- apply(rho, 1, which.max)
```

Notice that, `prior` is a scalar value, `mu` is a one dimensional array of length p , and `mu_matrix` is an array of size p by n which contains repeated columns of `mu`. In addition, `rho` is an array containing the LDA scores for all samples for each class. For comparative purposes, this computation will be performed twice: once for each estimate of the precision matrix. Further details on the parameters used and the results attained in this experiment see [3].

References

- [1] BOLLHÖFER, M., EFTEKHARI, A., SCHEIDEGGER, S., AND SCHENK, O. Large-scale Sparse Inverse Covariance Matrix Estimation. *SIAM Journal on Scientific Computing* 41, 1 (2019), A380–A401. URL: <https://doi.org/10.1137/17M1147615>.
- [2] EFTEKHARI, A., BOLLHÖFER, M., AND SCHENK, O. Distributed Memory Sparse Inverse Covariance Matrix Estimation on High-Performance Computing Architectures. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage, and Analysis* (2018), SC '18, IEEE Press. URL: <http://doi.org/10.1109/SC.2018.00023>.
- [3] EFTEKHARI, A., GAEDKE-MERZHÄUSER, L., PASADAKIS, D., BOLLHÖFER, M., SCHEIDEGGER, S., AND SCHENK, O. Sparse precision matrix estimation with SQUIC. *Under Review* (2023).
- [4] EFTEKHARI, A., PASADAKIS, D., BOLLHÖFER, M., SCHEIDEGGER, S., AND SCHENK, O. Block-enhanced precision matrix estimation for large-scale datasets. *Journal of Computational Science* (2021), 101389. URL: <https://doi.org/10.1016/j.jocs.2021.101389>.
- [5] HSIEH, C.-J., DHILLON, I., RAVIKUMAR, P., AND SUSTIK, M. Sparse inverse covariance matrix estimation using quadratic approximation. In *Advances in Neural Information Processing Systems* (2011), J. Shawe-Taylor, R. Zemel, P. Bartlett, F. Pereira, and K. Q. Weinberger, Eds., vol. 24, Curran Associates, Inc. URL: <https://proceedings.neurips.cc/paper/2011/hash/2ba8698b79439589fdd2b0f7218d8b07-Abstract.html>.